



A

ATOMICITY

C

CONSISTENCY

I

ISOLATION

D

DURABILITY

BASICALLY  
AVAILABLE

SOFT  
STATE

BASE

EVENTUALLY  
CONSISTENT

EVENTUAL  
CONSISTENCY

# CAP - Theorem

...or Brewer's theorem

- **Consistency**

Every client views the same data. Every node in distributed cluster receives the same data.

- **Availability**


During node failure time, every node must respond in a reasonable amount of time.

- **Partition tolerance**

Guarantees partition tolerance can recover from partitions once the partition heals.

The theorem is...

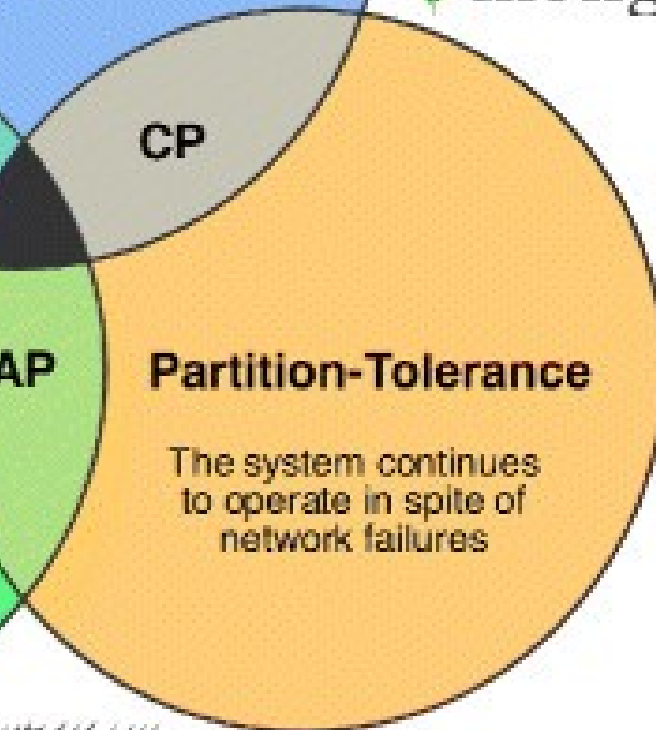

**You may only pick two!**



APACHE  
**HBASE**

mongoDB

redis

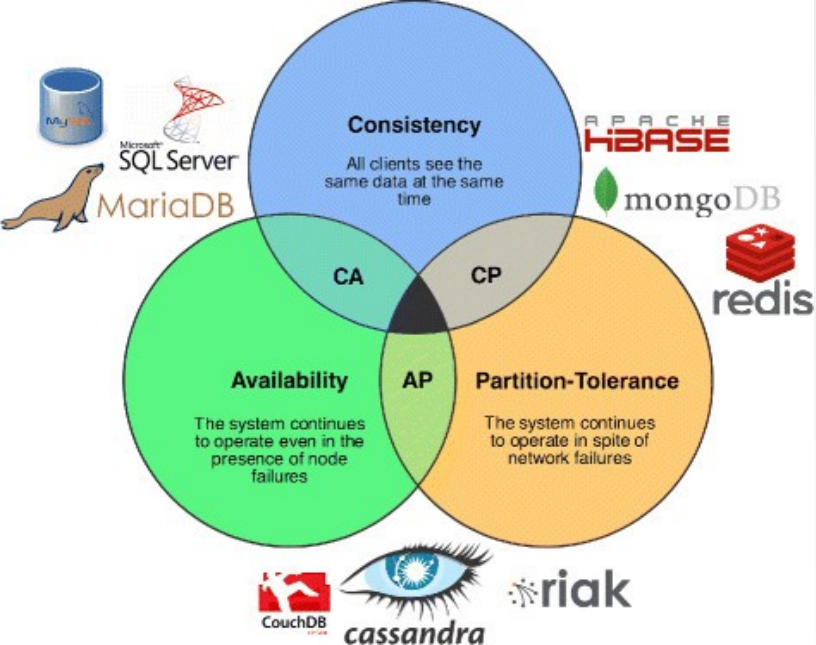


CouchDB

cassandra

riak






**AP (Availability and Partition tolerance):** When availability is chosen over consistency, the system will always process the client request and try to return the most recent available version of the information even if it cannot guarantee it is up to date due to network partitioning.

**CA (Consistency and Availability):** The correct value is delivered at all times, but no resilience to partitioning at all. Kafka is CA.

**CP (Consistency and Partition tolerance):** When consistency is chosen over availability, the system will return an error or time-out if particular information cannot be updated to other nodes due to network partition or failures.



ACID




APACHE  
**HBASE**

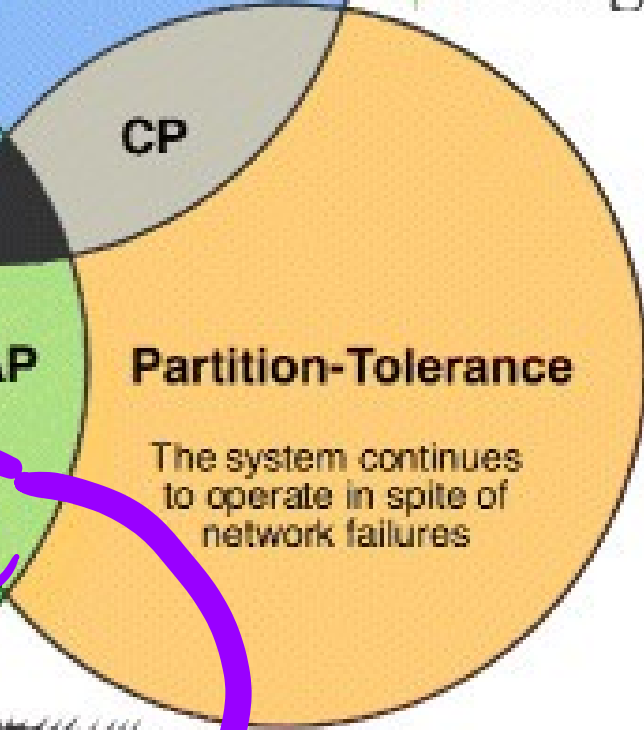
 mongoDB



redis



BASE



Couchbase



 riak

casandra

# ACID

**A**tomicity

**C**onsistency

**I**solation

**D**urability



# ACID

In 1983, [Andreas Reuter](#) and [Theo Härder](#) coined the acronym *ACID*, building on earlier work by [Jim Gray](#).

It describes the four properties that are the major guarantees of the transaction paradigm.

According to Gray and Reuter, the [IBM Information Management System](#) supported ACID transactions as early as 1973 (although the acronym was created later).



The diagram illustrates the ACID properties of databases. It consists of four vertical rectangular panels, each with a different background color: light green, medium green, blue, and purple. Each panel contains a white circle at the top with a red italicized letter inside. Below each circle is the name of the property in bold white text, followed by a brief description in white text. A large, light green double-headed arrow is positioned at the bottom, spanning the width of all four panels.

*a*

**Atomicity:**  
Transactions  
are all or  
nothing

*c*

**Consistency:**  
Only valid data  
is saved

*i*

**Isolation:**  
Transactions  
do not affect  
each other

*d*

**Durability:**  
Written data  
will not be lost



# Atomicity

An atomic system must guarantee atomicity in each and every situation, including power failures, errors, and crashes.

Transactions consist of statements and atomicity guarantees that a transaction is either fully committed or not committed at all.

# Consistency

Consistency ensures that a transaction can only bring the database from one valid state to another.

# Isolation

Transactions are often executed concurrently (e.g., multiple transactions reading and writing to a table at the same time).

Isolation ensures that concurrent execution of transactions leaves the database in the same state that would have been obtained if the transactions were executed sequentially.

# Isolation

Isolation is the main goal of concurrency control; depending on the method used, the effects of an incomplete transaction might not even be visible to other transactions.

# Durability

Durability guarantees that once a transaction has been committed, it will remain committed even in the case of a system failure (e.g., power outage or crash).

This usually means that completed transactions (or their effects) are recorded in non-volatile memory.

# BASE


**B**asic **A**vailability

**S**oft-state


**E**ventually consistent

## NoSQL


Key-Value



Column-Family



Graph



Document



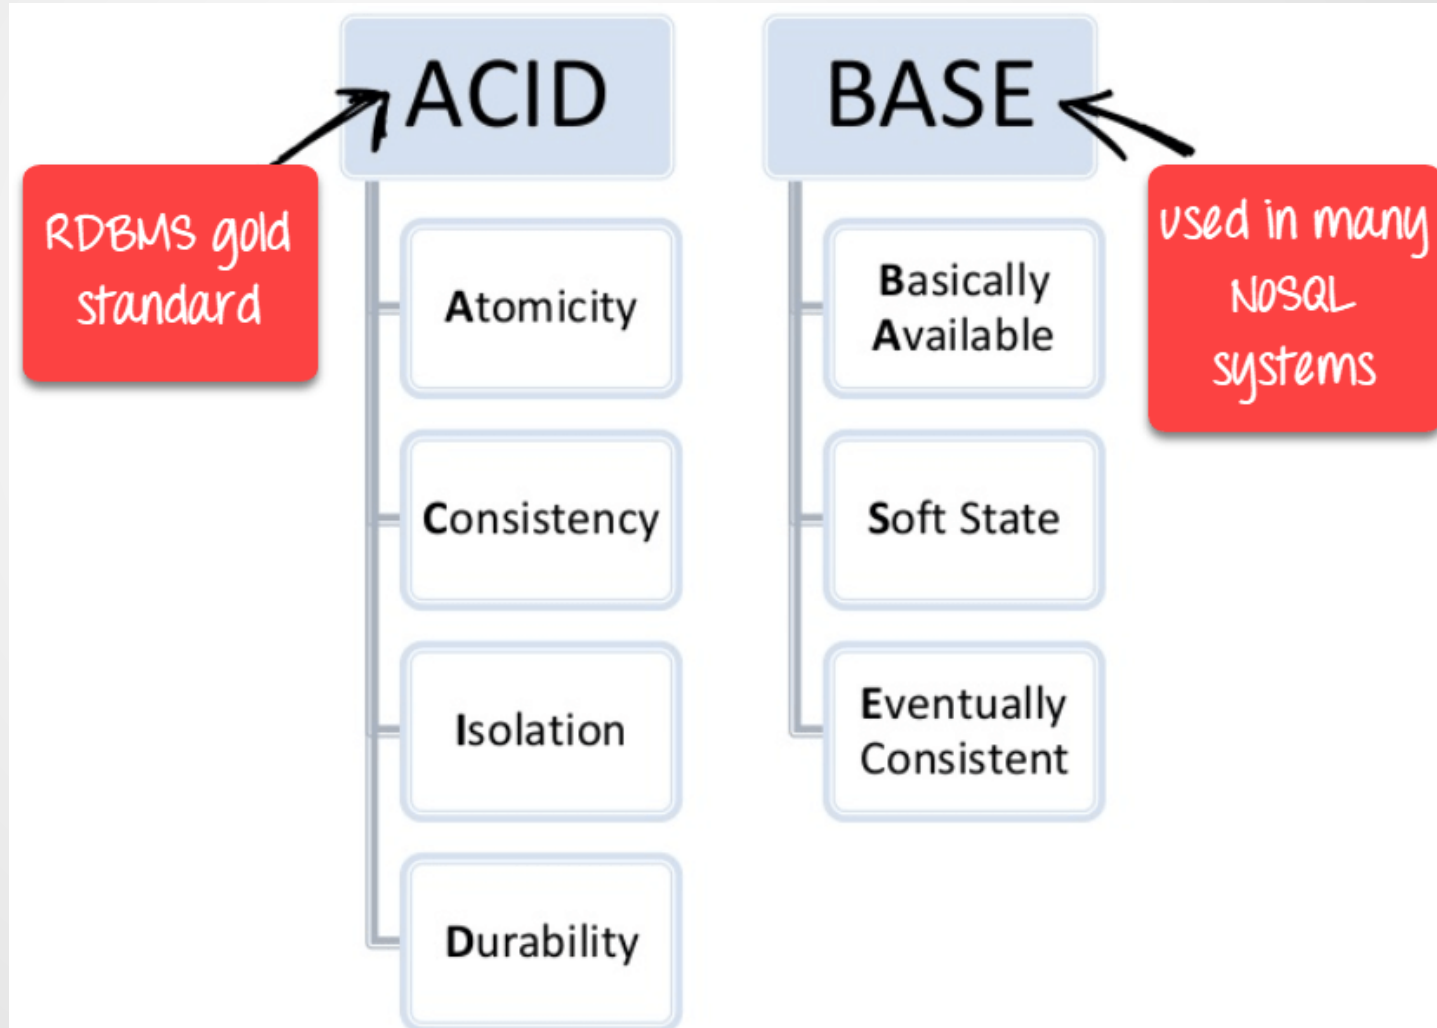
*DynamoDB,  
Cassandra,  
CouchDB,  
SimpleDB*

# Eventual Consistency

**Eventual consistency** is a consistency model used in distributed computing to achieve high availability.

It informally guarantees that, if no new updates are made to a given data item, eventually all accesses to that item will return the last updated value.

# BASE





# Basically Available

Reading and writing operations are available as much as possible (using all nodes of a database cluster), but might not be consistent.

(The write might not persist after conflicts are reconciled, and the read might not get the latest write.)

# Soft-State

Without consistency guarantees, after some amount of time, we only have some probability of knowing the state, since it might not yet have converged.

# Eventually Consistent

If we execute some writes and then the system functions long enough, we can know the state of the data.



Any further reads of that data item will return the same value.

It is purely a liveness guarantee (reads eventually return the same value) and does not guarantee safety.

An eventually consistent system can return **any value** before it converges.

# Eventual Consistency

Example







# References

- <https://www.techopedia.com/definition/29165/eventual-consistency>
- <https://en.wikipedia.org/wiki/ACID>
- [https://en.wikipedia.org/wiki/Eventual\\_consistency](https://en.wikipedia.org/wiki/Eventual_consistency)
- <https://systemdesignbasic.wordpress.com/2020/06/20/15-acid-vs-base-database-transactions/>
- <https://www.linkedin.com/pulse/what-acid-properties-database-aseem-jain/>
- [https://en.wikipedia.org/wiki/CAP\\_theorem](https://en.wikipedia.org/wiki/CAP_theorem)
- [https://www.debadityachakravorty.com/system\\_design/captheorem/](https://www.debadityachakravorty.com/system_design/captheorem/)